

Midterm

Student name: *Zahin Mohammad*

z5mohamm

20669584

Course: *ECE 495* – Professor: *Krzysztof Czarnecki*

Due date: *March 3rd, 2021*

Q1. ADS Fundamentals [1 mark]

1a. A vehicle is equipped with a stop-and-go pilot, which can fully operate a vehicle on a highway in traffic jams, but with a fallback-ready user. Which level(s) of driving automation is this driving automation system operating at? [1 mark]

Answer. TODO

Q2. Computer Vision Fundamentals [11 marks]

2a. Compute 1-D cross-correlation by applying the following filter $[0 \ 2 \ 1]$ to the following signal $[0 \ 1 \ 3 \ 0]$ (assume enough zero padding to show all non-zero output). [2 marks]

Answer. TODO

2b. Assume that the output of cross-correlating the 1-D filter $[2 \ 3 \ 2]$ with some input signal resulted in the following output signal $[7 \ 10 \ 7]$. What would be the output signal have we used convolution instead of cross-correlation and why? [1 mark]

Answer. TODO

2c. What would be the result of applying this filter to an image (Hint: add the Gaussian kernel to this filter from question 2e to recognize it)? [1 mark] Explain.

$$\frac{1}{16} \begin{bmatrix} -1 & -2 & -1 \\ -2 & 28 & -2 \\ -1 & -2 & -1 \end{bmatrix}$$

Answer. TODO

2d. What is the name of the following filter and what is it computing? [1 mark]

$$\frac{1}{16} \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 1 & 0 \\ 0 & 2 & 4 & 2 & 0 \\ 0 & 1 & 2 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} - \frac{1}{273} \begin{bmatrix} 1 & 4 & 7 & 4 & 1 \\ 4 & 16 & 26 & 16 & 4 \\ 7 & 26 & 41 & 26 & 7 \\ 4 & 16 & 26 & 16 & 4 \\ 1 & 4 & 7 & 4 & 1 \end{bmatrix}$$

Answer. TODO

2e. Complete the Gaussian filter shown below by specifying the cells marked by "?". How did you determine the missing values? [2 marks]


$$\frac{1}{16} \begin{bmatrix} 1 & 2 & ? \\ ? & ? & ? \\ ? & ? & ? \end{bmatrix}$$

Answer. TODO

2f. Why is the Canny filter using double thresholding? [1 mark]

Answer. TODO

2g. Consider the following representation in Hough space (polar coordinates).



2g.1 What is each of the individual sinusoids corresponding to in the input image? [1 mark]

Answer. TODO

2g.2 What is each of the black dots (sinusoid intersections) corresponding to in the input image (note that the horizontal axis sweeps 360 degrees)? [1 mark]

Answer. TODO

2h. Which of the following value profiles represent black in HSV? (select all that apply) [1 mark]

1. 0,high,high
2. any,low,low
3. any,low,high
4. 60,high,high
5. any,high,low
6. any,any,low

Answer. TODO

Q3. Machine Learning Fundamentals [21 marks]

3a. Starting with the probabilistic model for linear regression (assume single input x and single out y), show that maximizing the likelihood for a dataset (x, y) (with i.i.d. datapoints) is equivalent to minimizing the sum of squared errors. Hint: go via negative log likelihood [5 marks]

Answer. TODO

3b. What is the regression loss (as used in class) for a data point with label 0.4 and predicted output 0.7? [1 mark]

Answer. TODO

3c. (Apply what you've learned in lecture and Assignment 2) Given the following set of input vector X , ground truth vector Y , and weight matrices W_1 , W_2 , B_1 , and B_2 of a 2 layer fully connected neural network, what is the inference probability of the correct class? What is the cross-entropy loss value? Assume ReLU activation on the first hidden layer and softmax activation on the output layer. Show each step of the computation. Hint: Use numerically stable softmax and assume $e^{-1} \approx 0.37$ and $e^{-7} \approx 0.00$ [5 marks]

X	W1		B1	W2			B2	Y
2	10	8	-25	3	1	1	-3	0 (Man)
0	-5	3	11	1	2	3	3	1 (Moose)
	4	2	-10	5	4	2	0	0 (Deer)

Answer. TODO

3d. Consider the computational graph below for the following function

$$f(x_1, x_2) = \ln(3x_1 + e^{2x_2}) \quad (1)$$

Draw the computational graph and annotate it with the forward pass (above the arrows) and backward pass (below the arrows) for $x_1 = 1$ and $x_2 = 0$ (propagate the gradient back to each function input). Recall

$$\frac{de^x}{dx} = e^x \quad (2)$$

$$\frac{d\ln(x)}{dx} = \frac{1}{x} \quad (3)$$

Assume $\ln(4) \approx 1.39$ [5 marks]

Answer. TODO

3e. What is the difference between Stochastic Gradient Descent and ordinary (Batch) Gradient Descent? [1 mark]

Answer. TODO

3f. How is the condition of overfitting defined? [1 mark]

Answer. TODO

3g. Consider a convolutional layer with an input volume of depth 4 and output volume of depth 128. How many convolutional filters does the layer contain? What is the depth of each filter? [2 marks]

Answer. TODO

Q4. Semantic Segmentation [2 marks]

4a. (Apply what you've learned in lecture and Assignment 3) Semantic segmentation architectures sometimes use skip connections from early feature maps of the feature extractor to the corresponding-size upsampled maps in the decoder. What is the role of these connections? [1 mark]

Answer. TODO

4b. (Assignment 3) You were recommended to use batch norm as part of your network. Given a layer with the linear transformation and an activation function, where is the batch norm operation normally applied? [1 mark]

Answer. TODO

Q5. Object Detection [9 marks]


5a. Assume that an object detector uses a 5-by-5 grid and 3 anchor boxes at each cell. What is the maximum number of objects that the detector can detect? [1 mark]

Answer. TODO

5b. On a test set with a total of 4 of cars in the ground truth, a detector produced 3 bounding boxes with the following (score, IoU): (0.5, 0.7), (0.7, 0.8), (0.9, 0.2) (assume that each returned bounding box overlaps with a different ground truth). Assuming a score threshold of 0.6 and IoU threshold of 0.6, specify the number of TP, FP, and FN. [3 marks]

Answer. TODO

5c. (Assignment 4) Consider the figure below, which shows the bounding boxes predicted by a network (before non-maximum suppression). With an IoU threshold of 0.5 and detection threshold of 0.2, how many bounding boxes are left after non-maximal suppression? Explain your reasoning. [3 marks]



Answer. TODO

5d. How is it possible for two different bounding boxes to be generated for the same object (before non-maximum suppression)? [1mark]

Answer. TODO

5e. Which of the statements is correct? [1mark]

- a) An output neuron in an object detector is influenced only by input pixels within the anchor box assigned to it.
- b) An output neuron in an object detector is influenced only by input pixels within its positive anchor box.
- c) An output neuron in an object detector is influenced only by input pixels within its empirical receptive field.

Answer. TODO